Enterprise Data Readiness for AI

Twelve sessions of two hours for leaders and their data teams. You leave with a scored baseline, a working search and a twelve-month plan.

Bring
A decision maker and someone from data or IT, from the same organisation. A laptop and 50 documents that are not confidential.

Module and sessions	What you do	Tools	What you make
Start with a baseline 1 – 2	Name the AI uses you want and the data each needs. Score your organisation on five dimensions, with evidence for each score.	Spreadsheet scorecard	A baseline scorecard and three use cases
Know your data 3 – 4	List every data source with an owner and a sensitivity level. Write quality checks and a data contract for one dataset.	DataHub, OpenMetadata, Pandera, Great Expectations Core	A data inventory and one data contract
Make it searchable 5 – 6 – 7	Turn PDFs and scans into clean text. Search by meaning, then test it on questions your staff really ask, with access rules on.	Docling, sentence-transformers, pgvector, Chroma, Ragas	A working search over your own documents, and its test results
Platform, law and cost 8 – 9 – 10	Design the path from source to search with a check at each step. Map your data to four national laws. Price a hosting option.	DuckDB, Dagster, Presidio, Keycloak, Superset	A reference architecture, a data protection map and a cost estimate
Govern and plan 11 – 12	Set roles, a policy and a risk register using NIST and ISO 42001. Sequence a twelve-month plan and present it as a board paper.	NIST AI RMF, ISO/IEC 42001	A twelve-month roadmap and a board presentation

● Project milestones: sessions 2, 7 and 12

Case. We work on Kitara Group, a fictional bank and insurer, then on your own organisation.

Every session. Case 10 minutes, teach 30, live demo 25, break 5, workshop 40, decide 10.

Everything is free. No paid tool, cloud account or graphics card. We read each tool's licence in session 10.

KLAi Professional Learning. Tutor: Jotham Martin Wambi.