Learner Workbook

Leading with AI. One page for each session. Write in it, or type in Word.

How to use this workbook

Each page has the day's question, your steps for the hands-on, the fields to fill in, two questions to think about and one decision.

Every session adds to your AI Opportunity Portfolio (the Templates folder has it as a spreadsheet). The templates named on a page are in your Templates folder.

Work on your own organisation from the start. The fictional Nile Valley Foods examples show what a filled entry looks like.

My industry and organisation (fill in at Session 0)

My industry	
My organisation and role	
Two decisions or workflows I want AI to help with	
AI tools my staff already use	
What success would look like in a year	

Session 0, the introduction

The AI landscape for leaders

What is AI, in plain words, and where is it already at work in my industry?

Your steps for today

- Circle your industry for the course.
- Write two entries in your AI Opportunity Portfolio, each for a different job.
- Ask your team the three questions before Session 1 and note the answers.
- Write your first step for the next two weeks.

Fill in

My industry for the course	
Portfolio entry 1: use, in, out, who uses it, how we would know it works, which job	
Portfolio entry 2: use, in, out, who uses it, how we would know it works, which job	
What my team told me about the AI tools they use	
My first step: ask, protect, pick or wait, and why	

Think about

Which tile of the industry tour is closest to my business, and which case surprised me?

Where in my organisation is a mistake cheap and easy to check?

My decision

My first step in the next two weeks, and the reason.	

I keep: Your chosen industry and the first page of your AI Opportunity Portfolio, with two entries. Session 1 adds a use where AI sees or reads for you.

Session 1

AI that sees and reads: images, documents and speech

What can AI see or read for us, and when should we not trust what it sees?

Your steps for today

- Find a pile of photos, paper or recordings that people in your organisation handle by eye or by hand.
- Write it as a portfolio entry: the use, what goes in, what comes out, who uses it, how we would know it works.
- Add the Session 1 column: what exactly it sees or reads, and what happens to a poor picture or an unclear recording.
- Note how your pictures differ from a vendor's: cameras, light, forms, handwriting, languages.
- Write one question for a vendor and name who would check what the tool flags.

Fill in

The pile and who handles it today	
What goes in	
What comes out	
Who uses it	
How we would know it works	
What it sees or reads	
What happens to a poor picture	
How our pictures differ	
Who checks what it flags	
One question for the vendor	

Think about

Which pile in my organisation is largest, and who checks it today?

Could we collect a hundred of our own samples to test a tool on?

My decision

Paper to data, photos to check, speech to text, or not yet, with the reason.	

I keep: A one-page sheet: one idea behind seeing, reading and hearing, when it struggles, the questions to ask, today's recorded results, and your new portfolio entry.

Session 2

AI that predicts from your data

Who will default, leave, buy, fall sick or break down next, and can we trust the answer?

Your steps for today

- Choose a decision your organisation makes often, with past records behind it.
- Write it as a portfolio entry: the use, what goes in, what comes out, who uses it, how we would know it works.
- Add the Session 2 column: the kind of prediction (yes or no, a number, a ranking) and the decision it changes.
- Run the four data checks: history, labels, enough cases, a future like the past. Note any that fail.
- Write who would own the decision and who would set the cut-off.

Fill in

The decision	
The outcome to predict	
What goes in	
What comes out	
Who uses it	
How we would know it works	
Kind of prediction, and the decision it changes	
Data checks: history, labels, enough cases, future like the past	
Owner of the decision; who sets the cut-off	

Think about

Which of my predictions would someone act on next month?

Do we record the outcome of this decision today? If not, who could start?

My decision

Approve or refuse, plan how much, who first, or not yet, with the reason.	

I keep: A one-page sheet: the three kinds of prediction, the four data checks, the questions to ask, today's recorded results, and your new portfolio entries.

Session 3

Is it any good? Accuracy claims and the cost of mistakes

How do I judge a vendor who says "95% accurate", and what will its mistakes cost us?

Your steps for today

- For each portfolio entry, name the false alarm and the miss.
- Put a rough shilling figure on each, as a range if needed.
- Say which mistake costs more, and which way you would turn the dial.
- Write how you would test the tool on your own recent cases, and how you would check the mistakes by group.
- Name who would review the alarms and who would answer to a customer wrongly flagged.

Fill in

Portfolio entry	
A false alarm here means	
Cost of a false alarm (UGX)	
A miss here means	
Cost of a miss (UGX)	
Which costs more, and the dial setting	
How we would test it on our own cases	
Who reviews alarms; who answers to the customer	

Think about

Which of my entries has the most expensive miss?

How many false alarms a week could my team handle?

My decision

Catch more, fewer alarms, balanced, or not yet, with the cost that decided it.	

I keep: The vendor question sheet on accuracy, and a cost-of-a-mistake line for each entry in your AI Opportunity Portfolio.

Session 4

AI that understands language: sort, search, translate and summarise

What can AI do with our messages, contracts and calls, in English and in our local languages?

Your steps for today

- Name a pile of text your organisation struggles to read: messages, contracts, complaints, reports, minutes or recorded calls.
- Write it as a portfolio entry: the use, what goes in, what comes out, who uses it, how we would know it works.
- Add the Session 4 column: the job (sort, search by meaning, translate or summarise) and the languages involved.
- Write how a person would check the output each month, and how a customer would reach a person.
- Note any question from the six that your team or vendor cannot yet answer.

Fill in

The pile of text	
What goes in	
What comes out	
Who uses it	
How we would know it works	
The job: sort, search, translate or summarise	
Languages involved	
How a person checks it	
Open questions for the vendor	

Think about

Which pile of text costs us the most staff time to read today?

In which languages do our customers write to us, and who could test a tool in them?

My decision

Sort the inbox, search our documents, translate or summarise, or not yet, with the reason.	

I keep: A one-page sheet: the four language jobs, the map of meaning, today's recorded search results, the questions to ask, and your new portfolio entries.

Session 5

Generative AI: how chatbots write, and why they sometimes invent

Why does a chatbot sound so sure, and when can I trust what it writes?

Your steps for today

- Choose a writing task your staff repeat every week: replies, reports, minutes, letters.
- Write it as a portfolio entry: the use, what goes in, what comes out, who uses it, how we would know it works.
- Add the Session 5 column: what a person must check, and who signs.
- Mark which invention risks apply: citations, figures, policies, names.
- Write where the data you would paste in comes from, and whether it is confidential.

Fill in

The drafting task	
What goes in	
What comes out	
Who uses it	
How we would know it works	
What a person must check	
Who signs	
Invention risks: citations, figures, policies, names	
Data pasted in: source and sensitivity	

Think about

Which of my staff already use a chatbot for this task, with or without permission?

What would a wrong answer cost us, and who would notice first?

My decision

Draft replies, sum up reports, staff questions, or not yet, with the reason and the person who signs.	

I keep: A one-page sheet: how a chatbot writes, why answers vary, where it invents, the five questions to ask of any answer, today's recorded runs, and your new portfolio entry.

Session 6

The AI market: models, vendors, and buy, rent or build

Which model and supplier, where does our data go, and what will it cost to run?

Your steps for today

- List the entries in your AI Opportunity Portfolio from Sessions 1 to 5.
- For each, note the data it touches: public, customer, staff or health records.
- Estimate the monthly volume of use, even roughly.
- Add the Session 6 column: the route (rent, host, build on a partner), the reason, and the exit plan.
- Write the vendor questions you would ask first for your leading entry.

Fill in

The use	
What goes in	
What comes out	
Who uses it	
How we would know it works	
Data it touches	
Monthly volume	
Route and reason	
How we would leave	
Who owns the vendor relationship	

Think about

Which of our current AI tools could we leave in a month, and which could we not?

Do we know where each of our AI suppliers keeps our data?

My decision

Rent, host, build on a partner, or wait, with the one thing that would change your mind.	

I keep: A one-page sheet: the three routes, the four questions, the vendor questions, today's recorded comparison, and the route for each portfolio entry.

Session 7

Getting good work from AI assistants: prompting for leaders and teams

How do my staff get reliable work from these tools?

Your steps for today

- Choose one task your team does every week.
- Write its brief in the Reusable Brief Template: role, context, examples, format, constraints.
- Read it aloud as a new colleague and fix what you would still have to guess.
- List ten past cases whose right answer you know; run the brief on them this week and record the count in the Reliability Test Record.
- Add the Session 7 column to your AI Opportunity Portfolio: the tested brief, its count, and who checks.

Fill in

The use	
What goes in	
What comes out	
Who uses it	
How we would know it works	
Role	
Context	
Examples	
Format	
Constraints	
Ten test cases and their right answers	
Count right out of ten	
Who checks before it is sent	

Think about

Which of our tasks might sit outside the frontier?

Who on my team already uses an assistant, and what could the rest learn from them?

My decision

Drafting, summarising, sorting, or not yet, with the task, the count target out of ten, and who checks.	

I keep: A one-page sheet: the five parts of a brief, the test on ten tasks, the jagged frontier, today's recorded runs, and your team task with its brief, kept in the Reusable Brief Template and the Reliability Test Record.

Session 8

Making AI know your business: assistants on your own documents

How does AI answer from our policies, products and files, and how far can we trust what it says?

Your steps for today

- Pick a question your staff or customers ask again and again.
- Name the documents that answer it, and where each one lives today.
- Write it as a portfolio entry: the use, what goes in, what comes out, who uses it, how we would know it works.
- Add the Session 8 column: the route (prompt, retrieve or customise) and the documents it reads.
- Name the owner of each document, and mark any that is out of date, duplicated or missing.

Fill in

The question asked often	
The documents that answer it	
What goes in	
What comes out	
Who uses it	
How we would know it works	
Route: prompt, retrieve or customise	
Document owners	
Documents out of date, duplicated or missing	

Think about

Which of our documents would I be embarrassed for an assistant to quote?

Who would notice first if an answer came from an old version?

My decision

Prompt, retrieve, customise, or fix the documents first, with the reason.	

I keep: A one-page sheet: the open-book picture, the climb from prompt to retrieve to customise, the document checks, the questions to ask, today's recorded scores, and your new portfolio entry.

Session 9

When AI goes wrong: failures, risk and checking

What can go wrong with AI, what did it cost others, and how do we check before we trust it?

Your steps for today

- List the entries in your AI Opportunity Portfolio from Sessions 1 to 8.
- For each, name the worst failure, using the five: invents, misses the source, repeats bias, obeys hidden instructions, drifts.
- Price it: how likely (rare, sometimes, often) and how costly one error would be, in shillings.
- Choose the check: a ten-question test set, the five steps on each answer, or both, and name its owner.
- For your first entry, write five test questions with agreed answers, one of them not answered by your documents.

Fill in

Portfolio entry	
Worst failure (one of the five)	
How likely	
Cost of one error, UGX	
The check	
Owner of the check	
Five test questions and agreed answers	
Who compares results by group	

Think about

Which of our AI uses would I least like to explain to a regulator or a journalist?

How would a member of staff report an AI error today, and would they?

My decision

Go with a light check, go with a test set, a person signs, or not yet, with the reason.	

I keep: A one-page sheet: the five failures, the five-step check, the test-set template, the cost of an error in shillings, the questions to ask, and your portfolio entries with their worst failure.

Session 10

AI agents: delegating work that acts

What may an AI agent do on its own, and who approves?

Your steps for today

- For every entry in your AI Opportunity Portfolio, write the Session 10 column: the level of freedom from 0 to 4 (answers, suggests, drafts, acts then reports, acts alone).
- Choose one entry that would act. Fill a Tool Boundary Sheet for it: for read, write, send and spend, what it may do, what a person approves and what it must never do.
- Name the owner, where the log is kept, the stop rules and who runs the off switch.
- Write the one key you would never give it, and why.

Fill in

The entry	
Level of freedom (0 to 4)	
Keys it needs: read, write, send, spend	
What a person approves	
What it must never do	
Owner; where the log is kept	
Stop rules	
Off switch: who, and how fast	

Think about

Which tool already in use in my organisation acts without a person approving?

Who would notice first if an agent here went wrong, and how?

My decision

Level 1, level 2, level 3 or not yet, for my first agent, with the reason.	

I keep: A one-page sheet: the agent loop, the four kinds of key, the five levels of freedom, the controls, the signs of a hidden instruction, the questions to ask, and your portfolio with a level for each entry.

Session 11

Redesigning work: workflows, value and the business case (practical)

Which tasks change, what is it worth, and when does it pay back?

Your steps for today

- Choose one task your team repeats every week. Fill the top of a Task Card: task, owner, how often, time today.
- Map it in lanes: trigger, steps, handoffs (red) and the check (gold).
- For each step write a rule, an AI step, an agent (with its level of freedom) or a person, with one reason and how a person checks it.
- Fill the Inputs sheet of the Business Case Workbook for the step that costs the most time. Mark each figure E (estimate) or M (measured). Read the net gain a month, the payback period and the first-year gain, before and after the haircut.
- Score every entry of your AI Opportunity Portfolio on the Score sheet: value 1 to 5, feasibility 1 to 5. Circle the top entry that is cheap to get wrong.
- Write the four numbers you will record for four weeks, and your stop rule.

Fill in

My workflow and its trigger	
The handoff where work waits longest	
The step I would change first, and the kind of help	
Baseline: minutes each time today, times a month	
Net gain a month (UGX), payback period (months), first-year gain (UGX), after the haircut (UGX)	
Top portfolio entry: value, feasibility, priority	
Roles the change touches, and the hours that move	
Four numbers to record; my stop rule	

Think about

Which step on my map is a rule I could set up this month without AI?

Which of my figures did I measure, and which did I guess?

My decision

A rule first, one AI step, an agent at level 2, or measure first, for my top entry, with its payback.	

I keep: Your mapped workflow with a choice on every step, your Business Case Workbook with the payback and the first-year gain, and your portfolio scored on value and feasibility.

Session 12

Leading people through AI: roles, skills, adoption and change

What happens to my employees when AI arrives, and how do I bring them with me?

Your steps for today

- Take the portfolio entry you ranked first in Session 11 and the role it touches most.
- People and Skills Plan, part 1: list every task in that role today; for each, the AI's part, time on it (less, more, same, new) and the skill it needs.
- Part 2: what everyone, power users, builders and leaders learn, and how.
- Part 3: the three fears (job, being watched, being blamed): what you will say and what you will do. Cross out any promise you cannot keep.
- Part 4: the why in one sentence, what is approved and not approved, two champions, the feedback route, check-ins at day 30, 60 and 90.
- Write the one-page employee briefing and read it aloud in under two minutes.
- AI Opportunity Portfolio, sheet People, risk, data: roles it touches and tasks that change, for every entry.

Fill in

AI use, sponsor and plan owner	
The role, task by task: task today, with AI, time on it, skill it needs	
Skills by level: everyone, power users, builders, leaders	
Fears and answers: what we say, what we do	
The why, in one sentence	
Approved and not approved	
Champions, feedback route and check-ins	
Portfolio: roles each entry touches, tasks that change	

Think about

Which promise to my staff would be hardest to keep?

Who carries the change in my organisation: staff, contractors or new hires?

Who would my people trust as champions?

My decision

Train first, volunteers, redesign roles or not yet: my route, the reason and the date of the briefing.	

I keep: Your People and Skills Plan for your first AI use, the employee briefing, and the people column of your AI Opportunity Portfolio.

Session 13

Governance and buying: policy, risk register and vendors

What rules, risk register and vendor checks does my organisation need for AI?

Your steps for today

- Risk Register: edit the eight grey rows for your own AI uses; score likelihood and impact from 1 to 5; an owner and a control for your top five; add one risk of your own.
- AI Use Policy Template: draft sections 3 (approved tools), 4 (data), 5 (human review), 7 (accounts and owners), 8 (incidents) and 10 (review date).
- Lawyer page: mark which uses involve personal data, data sent abroad or EU customers; write the date you will see your lawyer.
- Vendor Question List for one tool: answer the data, security, and contract and exit groups; five questions with red-flag answers; the first-year cost in shillings.
- AI Opportunity Portfolio, sheet People, risk, data: likelihood and impact for every entry.
- Board briefing: one page from today's drafts, ending with what you ask the board.

Fill in

My top five risks: score, level, owner, control	
Approved tools and their owners	
What may go in at each data level	
What always needs a named reviewer	
Incident contact and how fast	
Policy owner, deputy and review date	
My uses with personal data; data sent abroad	
Five vendor questions and the red-flag answer for each	
First-year cost in shillings, every line	
Portfolio: risk score for each entry	
Board briefing: what we ask the board	

Think about

Which risk surprised me when I scored it?

Who in my organisation already uses AI through a personal account?

Which vendor claim would I have believed before today?

My decision

Decline, ask in writing or run a pilot: my route for the next proposal, the reason and the first step.	

I keep: Your AI Use Policy draft, your Risk Register, your lawyer page, your Vendor Question List, the risk column of your portfolio, and a board briefing outline.

Session 14

Data readiness, AI roadmap and your 90-day plan

Is our data ready, and what do we do in the next 90 days?

Your steps for today

- Build the inventory in the Information Inventory: at least six sources, each with an owner, a sensitivity label and a Ready for AI verdict.
- Rate the six areas Green, Amber or Red with one line of evidence each, read the total out of 12 and pick the top three gaps.
- Finish the AI Opportunity Portfolio (sheet People, risk, data): data readiness, a decision (Start now, Prepare first, Later, Drop) and one owner per entry; one entry to start now.
- Turn the gaps and the start-now entry into at least six roadmap lines with owner, dates, hours, what each waits for and a success measure.
- Put the day 30, 60 and 90 check-ins in the diary, list three risks to the plan and the decisions it needs, and write the day 90 target.

Fill in

My organisation and my plan start date	
Sources in my inventory, and how many are Confidential or Personal	
My six ratings, my total out of 12, my first Red	
My top three gaps, with evidence and a first fix	
My final portfolio: the entry I start now, those I prepare first, those I drop	
My roadmap lines: action, owner, due date, hours, waits for, success measure	
My sponsor and my data owners	
Three risks to the plan and what I will do	
Decisions needed, who decides, by when	
My check-in dates: day 30, day 60, day 90	
My day 90 target: the score and one count	
The action I take within 24 hours	

Think about

Which area surprised me when I scored it?

Which portfolio entry did I drop, and why?

Which of the fifteen sessions changed how I think most?

What is AI, and where is it at work in my industry? My answer in two sentences.

My decision

How much my plan delivers by day 90 (one use finished, three gaps started, one folder tested, or data first), and my first action within 24 hours.	

I keep: Your inventory and scorecard, your final AI Opportunity Portfolio, your 90-day roadmap with owners and check-ins, and the Follow-up Kit.

Words used in this course

Word	What it means	First used
AI	Artificial intelligence: software that does tasks we link with human thinking, such as reading, predicting and writing.	Session 0
Machine learning	Software that learns a pattern from past examples, where a rule would be written by a person.	Session 0
Model	The pattern that machine learning produces, used to make a guess about a new case.	Session 0
Generative AI	Machine learning that writes new text, pictures or code; ChatGPT is one product of this kind.	Session 0
AI Opportunity Portfolio	Your one-page register of possible AI uses in your organisation. It grows every session.	Session 0
Pixel	One dot of a picture. The computer reads each dot as a number.	Session 1
Training data	The labelled examples a model learns from, such as photos marked healthy or diseased.	Session 1
Neural network	A model built in layers that finds small details, then shapes, then the answer.	Session 1
OCR	Optical character recognition: software that reads printed or handwritten letters from a scan or photo.	Session 1
Speech recognition	Software that turns recorded speech into written text, also called speech-to-text.	Session 1
Label	The known outcome in a past record, such as repaid or not repaid.	Session 2
Score	The model's estimate of how likely an outcome is, such as 8 in 10 likely to repay.	Session 2
Threshold	The cut-off score your policy sets: above it one action, below it another.	Session 2
Test set	Past records kept aside and never shown to the model, used to check it fairly.	Session 2
Overfitting	Learning the slips and accidents of past cases, so the model does worse on new ones.	Session 2
False positive	A false alarm: the system flags a case that was fine.	Session 3
False negative	A miss: the system lets through a case it should have caught.	Session 3
Accuracy	The share of all cases a system gets right. It hides both mistakes when the event is rare.	Session 3
Bias	When a system's mistakes fall more often on one group of people.	Session 3
Classification	Sorting each message into one of a few labels, such as order, payment or complaint.	Session 4
Semantic search	Search by meaning: finding the passage that means the same thing, even in different words.	Session 4
Embedding	A text's position on a map of meaning, worked out by a model. Close positions mean close meanings.	Session 4
Machine translation	AI that turns text in one language into text in another.	Session 4
Low-resource language	A language with little written text online for AI to learn from, as with most Ugandan languages.	Session 4
Large language model	The engine inside a chatbot: a model trained on a huge library of text to choose the next word.	Session 5
Token	A piece of text the model reads and writes: a whole word or part of one.	Session 5

Word	What it means	First used
Temperature	A setting for how adventurous the choice of next word is: low plays safe, high takes long shots.	Session 5
Hallucination	The industry's name for a fluent answer that is false or invented, such as a court case that does not exist.	Session 5
Hosted model	A model that runs on the vendor's computers. You send your text over the internet and pay to use it.	Session 6
Open weights	A model whose learned numbers can be downloaded and run on computers you choose, under a licence.	Session 6
Fine-tuning	Further training of an existing model on your own examples, so it fits your language or your task.	Session 6
Lock-in	When leaving a supplier would cost so much time and money that you are stuck with it.	Session 6
Benchmark	A fixed public test with an answer key, used to compare models. It is someone else's test.	Session 6
Prompt	Everything you give an assistant: the request and what it needs to do it well. Think of it as a brief.	Session 7
Few-shot examples	One or two finished examples pasted into the prompt, so the assistant copies their pattern.	Session 7
Constraints	The limits a reply must keep: what it must not do, and what to say when it is unsure.	Session 7
Reasoning model	A model that writes out its working before it answers: slower, often better on many-step work.	Session 7
Jagged frontier	The uneven edge of what AI does well: excellent at some tasks, worse than staff at others that look similar.	Session 7
Knowledge base	The collection of your own documents an assistant may answer from: policies, product sheets, past reports.	Session 8
Retrieval	Finding the pages closest in meaning to a question, as search by meaning did in Session 4.	Session 8
RAG	Retrieval-augmented generation: find the right pages first, then write the answer from them.	Session 8
Context window	How much text an assistant can read at once, your question and pages included.	Session 8
Retrieval miss	The search fetches the wrong page, an old page, or none, so the answer rests on the wrong source.	Session 9
Prompt injection	Instructions hidden in an email or document that try to make the AI do something you never asked.	Session 9
Drift	A tool that worked last month does worse now, because the documents, the world or the model changed.	Session 9
Agent	A chatbot given tools and a loop, so it can take steps and act, and choose its own next step.	Session 10
Tool	An action a program carries out for the agent, such as read a list, save a draft, send or pay.	Session 10
Autonomy level	How much the agent may do before a person says yes: its level of freedom.	Session 10
Least privilege	Give the agent only the keys, or permissions, that the task needs.	Session 10
Human-in-the-loop	A person approves, edits or rejects an action before it runs.	Session 10

Word	What it means	First used
Workflow map	A drawing of a task that repeats: one lane per person, a box per step, arrows for the order and the handoffs.	Session 11
Rule-based automation	A machine follows fixed rules with no judgement, such as a bank standing order.	Session 11
AI step	One judgement by a model inside a fixed path, such as a spam filter, checked by a person.	Session 11
Baseline	How the task performs today, in time, cost and errors, measured before any change.	Session 11
Payback period	How many months until the savings repay what you spent at the start.	Session 11
Return on investment	What the first year gives back on top of what it cost, in shillings or as a percentage.	Session 11
Task	One piece of work inside a job, such as sorting messages or calling a customer back.	Session 12
Exposure	How far AI can affect a job's tasks. Exposed means affected; it does not mean lost.	Session 12
Augment and automate	Augment: AI helps and a person decides. Automate: AI does the task and a person checks.	Session 12
Adoption	How many people actually use a new tool in their daily work, and how often.	Session 12
Champion	A trusted colleague who tries the tool first and helps the team use it well.	Session 12
Reskilling	Training people for new tasks or a new role as their work changes.	Session 12
AI governance	The policies, roles and checks by which an organisation directs and controls its use of AI.	Session 13
Acceptable use policy	A written rule on who may use which AI tools, for what, and with what data.	Session 13
Risk register	A table of what could go wrong, each risk scored for likelihood and impact, with an owner and a control.	Session 13
Data classification	Sorting information into levels, such as public or confidential, each with its own handling rule.	Session 13
Vendor due diligence	Checking a supplier's claims, security and contract terms before you commit.	Session 13
Total cost of ownership	Everything you will pay over the life of a tool: set-up, fees, people's time and leaving.	Session 13
Data inventory	A written list of every data source you hold, with its owner, where it lives and who may see it.	Session 14
Data quality checks	Counts of what is missing, wrong, doubled or out of date. The industry calls them data quality dimensions.	Session 14
Data owner	The named person who decides who may use a source and signs off its fixes.	Session 14
Readiness scorecard	Six areas rated Green, Amber or Red by a written rule, scored out of 12.	Session 14
Roadmap and milestone	A plan of actions in order, with owners and dates. A milestone is a dated checkpoint.	Session 14
Success measure	A number someone else can check, such as 0 faults by 19 February. Also called a KPI, key performance indicator.	Session 14